

Max-Flow I and II Notes

Max-Flow I


Max-flow is a fundamental problem with many applications in logistics and routing etc.

1 Input

A flow network $G = (V, E, s, t, c)$ has:

- Network (i.e. directed graph) (V, E)
- One source node $s \in V$ and one target node $t \in V$
- Edge capacities $c : V^2 \rightarrow \mathbb{R}_{\geq 0}$ where $c(u, v) = 0$ if $(u, v) \notin E$

Intuition: Each edge is a one-way road, and the capacity is the number of lanes on that road.




Fundamental question: What is the max value of traffic from s to t that we can support?

1.1 Gross Flow


Gross flow is specified by a function $g : E \rightarrow \mathbb{R}_{\geq 0}$, satisfying the following properties.

- *Feasibility:* $0 \leq g(e) \leq c(e)$ for all $e \in E$.
- *Flow conservation:* $\sum_{(x,v) \in E} g(x, v) = \sum_{(v,y) \in E} g(v, y)$ for all $v \in V \setminus \{s, t\}$.

Example: Each edge is labelled by $g(e)/c(e)$ where $g(e)$ is the rate and $c(e)$ is the value. Here we can see all flows \leq capacities and “flow in = flow out” for every node.



However, this definition is not convenient when dealing with flow cycles of length 1 or 2:



1.2 Net flow

Firstly, we won't assume that s is a source and t is a sink (i.e. both can have in- and out- neighbors), so $c(u, v) = 0$ for all $(u, v) \notin E$. We now define *net flow*.

Net flow is specified by a function $f : V^2 \rightarrow \mathbb{R}$, satisfying the following properties:

- *Feasibility*: $f(u, v) \leq c(u, v)$ for all $u, v \in V$.
- *Flow conservation*: $\sum_{x \in V} f(v, x) = 0$ for all $v \in V \setminus \{s, t\}$.
- *Skew-symmetry*: $f(u, v) = -f(v, u)$ for all $u, v \in V$.


Few corollaries to note:

1. We can have $f(u, v) \leq 0$ when $(u, v) \notin E$.
2. $f(v, v) = 0$ for all $v \in V$.
3. If $(u, v), (v, u) \notin E$ then $f(u, v) = 0$.
4. $\sum_{f(x,v)>0} f(x, v) = \sum_{f(v,y)>0} f(v, y)$ for all $v \in V \setminus \{s, t\}$.

We define the **value of flow** f to be $|f| = \sum_{v \in V} f(s, v)$.


The **Max-Flow Problem** is to find a flow with maximal value for a given flow network G .

Simplest examples of non-zero flow:




2 Flow Decomposition

An observation: Any flow can be decomposed into a collection of flow cycles and s - t flow paths:



can be decomposed into



Formally, denote the *support* $\text{supp}_f(G)$ as the subgraph of G with only edges with positive flow.

Flow Decomposition Theorem.

For any flow f with $|f| \geq 0$, $\text{supp}_f(G)$ can be decomposed into a collection of flow cycles and s - t flow paths.


Proof.

We induct on the number of edges in $\text{supp}_f(G)$. The base case is where the support has no edges (so $|f| = 0$), and indeed it can be decomposed into an (empty) set of flow cycles and s - t flow paths. Now assume $\text{supp}_{f'}(G)$ can always be decomposed into flow cycles and s - t flow paths when $|f'| < L$. Let there be a flow f with $|f| = L$.

Suppose that we can find a cycle C in $\text{supp}_f(G)$. Then we can “reduce that cycle” by subtracting all the flows on the cycle by $f_{\min}(C) = \min_{e \in C} f(e)$. Formally, we define a new flow f' on the graph:

$$f'(u, v) = \begin{cases} f(u, v) - f_{\min}(C) & (u, v) \in C \\ -f(u, v) + f_{\min}(C) & (v, u) \in C \\ f(u, v) & \text{otherwise} \end{cases}$$

For example,



We can check that f' still satisfies feasibility, flow conservation and skew-symmetry, so f' is a flow.

Now that $\text{supp}_{f'}(G)$ has at least one edge fewer than $\text{supp}_f(G)$, we can apply the inductive hypothesis, decomposing $\text{supp}_{f'}(G)$ and then adding the flow cycle back to get $\text{supp}_f(G)$.

What if there are no cycles? Note that each vertex, that is not s nor t , with an in-edge must have an out-edge by flow conservation, so if we perform DFS starting from s , we are going to reach either t or s again. The latter case forms a cycle, so we instead reach just t .

Now we have an s - t path P , we reduce it by subtracting the flow by $f_{\min}(P) = \min_{e \in P} f(e)$. In other

words, we define again a new flow

$$f'(u, v) = \begin{cases} f(u, v) - f_{\min}(P) & (u, v) \in P \\ -f(u, v) + f_{\min}(P) & (v, u) \in P \\ f(u, v) & \text{otherwise} \end{cases}$$

and check that it still satisfies it being a flow. Then at least one edge is removed, and we can use the inductive hypothesis similarly. \square

Exercise. There is always a flow decomposition with at most $|E|$ flow cycles and s - t flow paths.

Exercise*. If $|f| > 0$, then there must be at least one s - t flow path in $\text{supp}_f(G)$.

3 Cuts

Let f^* be a max flow of G (need not be unique), and denote $F^* = |f^*|$.

Warmup. How to check whether $F^* > 0$ using flow decomposition?


Let G^* be the subgraph of G with only edges with positive capacities. Note: $\text{supp}_f(G)$ is the subgraph with positive *flow* edges, and hence depends on f ; whereas G^* is the subgraph with positive *capacity* edges, and hence is independent of flow.

By flow decomposition, $F^* > 0$ if and only if there is an s - t path in G^* .

So how do we certify when there is no s - t path in G^* ? We use a *cut*.

Let $\bar{S} = \{v \in V \mid \text{exists } s\text{-}v \text{ path in } G^*\}$. Then when $F^* = 0$, $s \in \bar{S}$ but $t \in V \setminus \bar{S}$.

More generally, we define an s - t **cut** to be a cut $(S, V \setminus S)$ such that $s \in S$ and $t \in V \setminus S$.



We also define the *capacity of a cut* $c(S) = c(S, V \setminus S) = \sum_{u \in S} \sum_{v \in V \setminus S} c(u, v)$. In other words, $c(S)$ is the total capacity of edges leaving S . Then

$$F^* = 0 \iff \text{no } s\text{-}t \text{ path in } G^* \iff \text{exists } s\text{-}t \text{ cut } S \text{ with } c(S) = 0.$$

Therefore, s - t paths and s - t cuts are dual to each other!

Minimum s - t cut problem: Given G , find an s - t cut of minimum capacity.

Let $S^* = \underset{S \text{ } s\text{-}t \text{ cut}}{\text{argmin}} c(S)$ be a cut with minimum capacity. Hence $F^* = 0 \iff c(S^*) = 0$.

Given a flow f , we can also define a *flow across a cut* $f(S) = f(S, V \setminus S) = \sum_{u \in S} \sum_{v \in V \setminus S} f(u, v)$. By feasibility, $f(S) \leq c(S)$ for any s - t cut S .

Claim. For any two s - t cuts S, S' , we have $f(S) = f(S')$.

Proof. By flow decomposition, f is a collection of flow cycles and s - t flow paths. For any cut S ,

- Flow cycles contribute 0 to both $f(S)$ because of skew-symmetry and the fact that cycles will zig-zag into and out of S .
- s - t flow paths contribute to $f(S)$ the amount equal to the flow of the path itself, because paths will zig-zag into and out of S , cancelling out except for the last zig, which contributes the flow.

In both cases, $f(S)$ did not depend on S , so $f(S) = f(S')$. □


Therefore, for any s - t cut S , $f(S) = f(\{s\}) = |f|$.

Weak Duality (Maxflow \leq Mincut). $|f^*| = f(S) \leq c(S) \leq c(S^*)$.

4 Increasing Flow


Question: Given a flow f , how to increase its value or conclude that $|f| = F^*$?

Our first idea is to find another s - t path and push more flow along it. But this does not always work. Consider the following example:



There are no more s - t paths to push flow, yet $|f| = 1$ is not the max flow!

New idea: *Undo some existing flows.* We can think of (u, v) as being there with capacity 0, so we can push flow along $s \rightarrow u \rightarrow v \rightarrow t$, undoing the flow:



Now we have $|f| = 2$! How do we ensure that 2 is a maximum? We can use weak duality and choose an appropriate cut S . By choosing $S = \{s\}$, we get $2 = |f| \leq F^* \leq c(S^*) \leq c(\{s\}) = 2$, confirming maxflow.


5 Residual Network

In light of the previous example, we introduce a **residual network** to find pushable flows.

Given a flow f on a flow network G , we define the **residual network** $G_f = (V, E_f, s, t, c_f)$ where

- $c_f(u, v) = c(u, v) - f(u, v)$
- $(u, v) \in E_f \iff c_f(u, v) > 0$

For the previous example, here is how the original flow network and its residual network look like:




Notice that a non-edge $(u, v) \notin E$ can be in E_f if $f(v, u) > 0$. This is because

$$c_f(u, v) = c(u, v) - f(u, v) = 0 + f(v, u) > 0.$$

Therefore this definition allows us to undo flow!

Another example of flow vs residual:



Observation: If f, f' are flows in G, G_f respectively, then $f + f'$ is a flow in G .

So to improve f , we just need to find a non-zero flow in G_f .

What if there is no non-zero flow in G_f ? Then by setting \bar{S} as the set of vertices reachable from s in G_f , we have $c_f(\bar{S}) = 0$, which rearranges to

$$\sum_{u \in \bar{S}} \sum_{v \in V \setminus \bar{S}} (c(u, v) - f(u, v)) = 0 \implies c(\bar{S}) = f(\bar{S}).$$

Applying weak duality, $c(\bar{S}) = f(\bar{S}) = |f| \leq c(S^*) \leq c(\bar{S})$, so f is a max-flow and \bar{S} is a min-cut!

In conclusion, we have

$$|f| = F^* \iff \text{no } s\text{-}t \text{ path in } G_f$$

6 Towards an Algorithm

Define an *augmenting path* as a directed s - t path in G_f . We have shown that we should always push additional flow along such a path P up to the (residual) bottleneck capacity $c_f(P) = \min_{e \in P} c_f(e)$, increasing $|f|$ by $c_f(P)$. When there is no more augmenting path, we have also shown that f is a max-flow. We can write this as the FORD-FULKERSON algorithm:

Algorithm 1 FORD-FULKERSON (1956)

- 1: Start with zero flow
 - 2: **while** G_f has augmenting path, **do**
 - 3: Find augmenting path P in G_f via DFS $\triangleright O(|E|)$
 - 4: Augment flow by pushing $c_f(P)$ $\triangleright O(|E|)$
 - 5: **end while**
-

Runtime. If all capacities are integers in $[0, C]$ then

$$\text{no. of augmentations} \leq |f| \leq c(\{s\}) \leq |V| \cdot C$$

and thus the runtime is $O(|E| \cdot |f|) = O(|E| \cdot |V| \cdot C)$. This is *pseudopolynomial*: It is polynomial in the *numerical value* C , i.e. polynomial in the *unary* coding $(1, 11, 111, \dots)$ of the numbers instead of binary.

By this algorithm, we also have the flow integrality theorem.

Flow Integrality Theorem. If all capacities are integers, there exists an integral max-flow.

Note that if the capacities are irrational, this algorithm may potentially take infinite runtime!

Max-Flow II

7 Maxflow-Mincut Theorem

Maxflow-Mincut Theorem (Strong Duality). The following are equivalent:

- There is an s - t cut $(S, V \setminus S)$ such that $c(S) = f(S) = |f|$.
- f is a max flow.
- There is no s - t path in G_f .

Proof.

(1) \Rightarrow (2). Use weak duality: $|f| \leq |f^*| \leq c(S^*) \leq c(S) = f(S) = |f|$

(2) \Rightarrow (3). Contrapositive: If there is an s - t path in G_f , push the flow!

(3) \Rightarrow (1). Proven in section 5.

8 Other Algorithms

8.1 Max Bottleneck Path Algorithm (MBP)

Algorithm 2 MAX BOTTLENECK PATH

- 1: Start with zero flow
 - 2: **while** G_f has augmenting path, **do**
 - 3: Find augmenting path P in G_f that has maximum $c_f(P)$ ▷ $O(|E| \log |V|)$
 - 4: Augment flow by pushing $c_f(P)$ ▷ $O(|E|)$
 - 5: **end while**
-

Runtime. It turns out that the number of iterations is $O(|E| \log |f|)$, so the runtime is

$$O(m \log n) \cdot O(m \log |f|) = O(m^2 \log n \log(nC))$$

which is *weakly polynomial*.

8.2 Edmonds-Karp

Algorithm 3 EDMONDS-KARP

- 1: Start with zero flow
 - 2: **while** G_f has augmenting path, **do**
 - 3: Find augmenting path P in G_f with minimal number of edges
 - 4: Augment flow by pushing $c_f(P)$
 - 5: **end while**
-

Runtime. $O(m^2n)$. Polynomial (no proof given)

8.3 Later Work

KING-RAO TARJAN (1994): $O(mn \log_{m/n \log n} n)$

ORTIN (2013): $O(mn)$

GOLDBERG-RAO (1998): $O(\min(m^{3/2}, mn^{2/3}) \log(n^2/m) \log C)$

LEE-SIDFORD (2014): $O(m\sqrt{n} \log^{O(1)} n \log C)$

MADRY (2013): $O(m^{10/7} \log^{O(1)} n)$ for $C = 1$


CHEN-KYNG-LIU-PENG-GUTENBERG-SACHDEVA (1998): $O(m^{1+o(1)} \log C)$

9 Application

A *matching* in a graph G is a set of edges that do not share end points. The size $|M|$ of a matching M is the number of edges in M .

Max Bipartite Matching Problem: Given a bipartite graph G (a graph on vertices $L \cup R$ disjoint, with edges $E \subseteq L \times R$), output a matching of maximum size.

Solution. The method is to create a source s that connects to all vertices in L , and let all vertices in R to connect to a target t . Then we assign a capacity of 1 to all the edges.



1. By flow integrality, since all capacities are integers, there is a max flow with 0 or 1 values on every edge and FORD-FULKERSON will find it.
2. If the maxflow f^* has value F^* then F^* of the edges (s, x) get flow 1 and the rest get 0.
3. The max matching is $\geq F^*$ because given any maxflow f^* , no edges in $L \times R$ with flow 1 have a common node (otherwise the flow in/out is > 1 but the capacities are all 1), and hence we can find a matching with that number of edges, i.e. F^* . The best case is at least as good as this.
4. F^* is at least the max-matching because given any max-matching we can extend the matching to s and t and give flow 1 to all of them, giving a valid flow. The maxflow must be better than this.
5. Therefore $F^* = \text{max-matching}$.

Runtime via FORD-FULKERSON. $O(m|f|) = O(m|M|) = O(mn)$.

10 Optional: Running Time of MBP

Let f_i be the flow computed after iteration i , decomposable into $\leq |E(G_{f_i})|$ flow cycles and s - t flow paths.

Let f_i^* be the max flow in G_{f_i} after i -th iteration.

At the beginning, $|f_0^*| = F^*$.

(Exercise) $f^* = f_i + f_i^*$ for all i .

Since $|E(G_{f_i})| \leq 2m$, there is always an s - t flow path P with flow value $f_i^*(P) \geq \frac{|f_i^*|}{2m}$. For this path P in G_{f_i} , we must have $c_{f_i}(P) \geq f_i^*(P) \geq \frac{|f_i^*|}{2m}$.

Therefore after iteration $(i + 1)$, the residual flow

$$f_{i+1}^* \leq |f_i^*| - \frac{|f_i^*|}{2m}$$

which solves to

$$f_{i+1}^* \leq \left(1 - \frac{1}{2m}\right)^i F^*.$$

Consider $I = (4m \ln F^*) + 1$ iterations, then

$$\begin{aligned} |f_I^*| &\leq \left(1 - \frac{1}{2m}\right)^{4m \ln F^*} F^* \\ &< \left(\frac{1}{e}\right)^{2 \ln F^*} F^* \\ &= \frac{1}{(F^*)^2} F^* \leq 1 \end{aligned}$$

and hence $f_I^* = 0$. Therefore we need $\leq O(m \log F^*)$ iterations. □